

Kubernetes 101 for Python Programmers

Learn the ABCs of Kubernetes and how to get started on using managed containers for your python development on your local machine and also for production deployments on a cloud provider-managed Kubernetes cluster.

Sachin Agarwal <sachin@bigbitbus.com>

DRAFT v.05

Tutorial Agenda

1. Background

- Docker container runtime

2. Kubernetes Architecture & Components

3. Kubernetes for Managed Containers

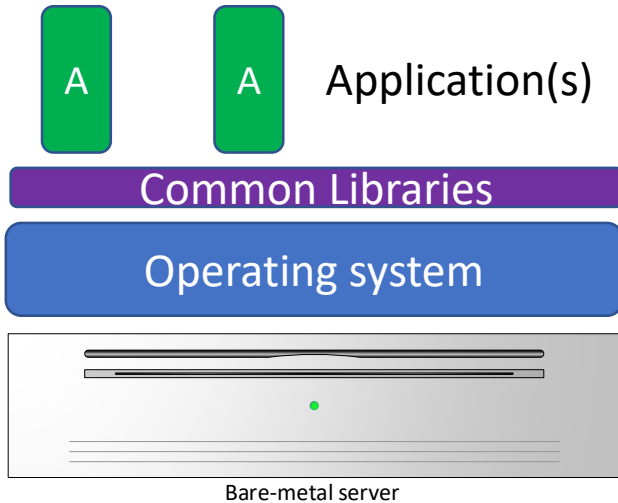
- Installation options – local laptop/cloud/DiY
- Services and deployments
- Operational constructs – scaling, updating code
- State considerations, secrets and environment variables

4. Further reading pointers

Bare-metal, Virtual Machines, Containers

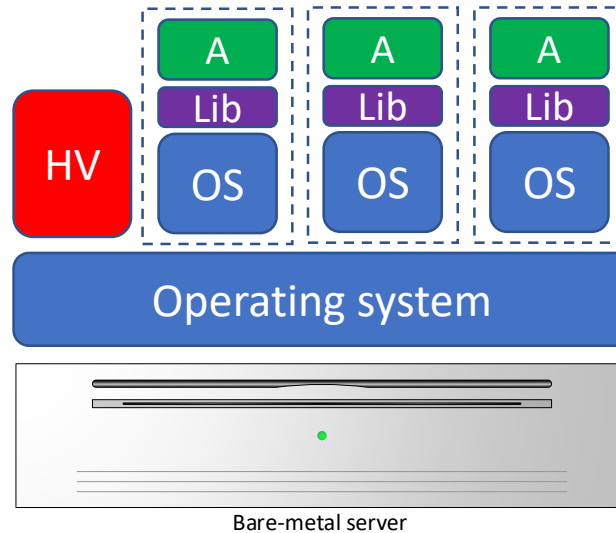
Bare Metal

- Weak encapsulation
- Long install/start-up times
- Ideal for single application e.g. a Hadoop node
- Hard to right-size




Virtual machines

- Strong encapsulation
- Hardware emulation overhead (per VM)
- Per-VM operating system overhead
- Long install/start-up times



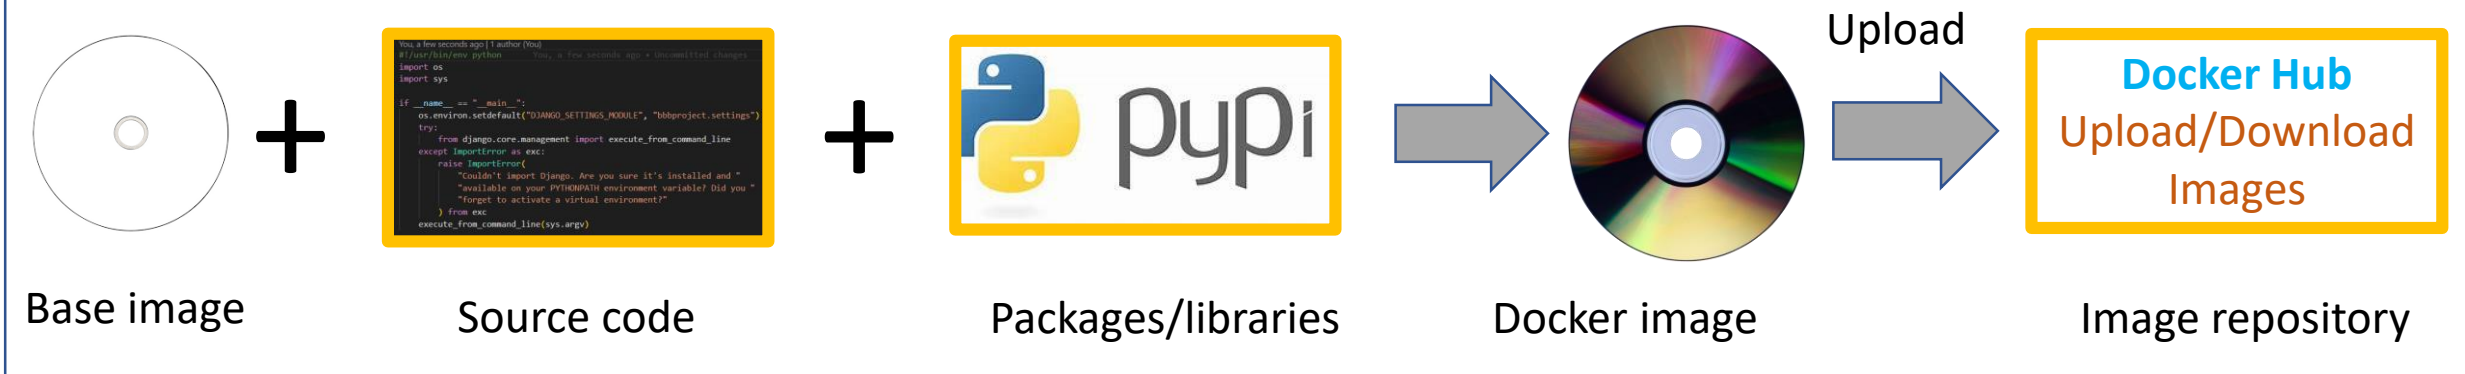
Containers

- Encapsulation *lite*
- Small overhead – only single kernel running
- Smaller deployment artifacts expressed as code
- Quick start-up times



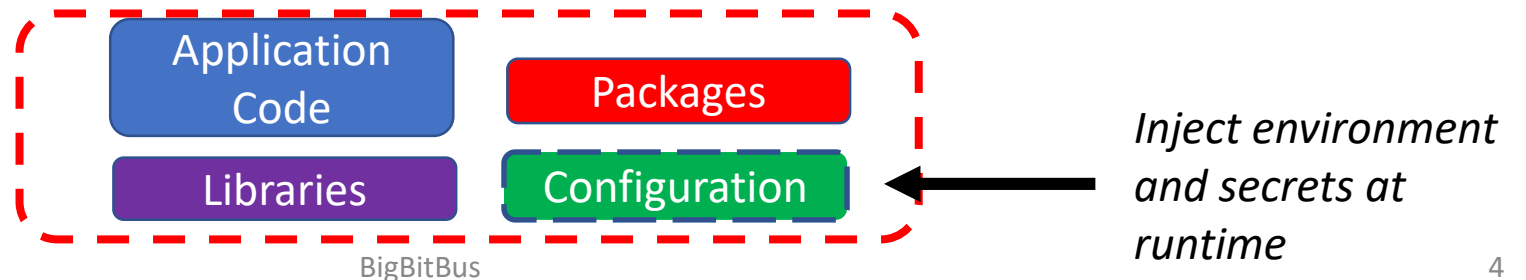
Docker – container runtime

Build




Run

- *Download immutable image*
- *Configure and run*

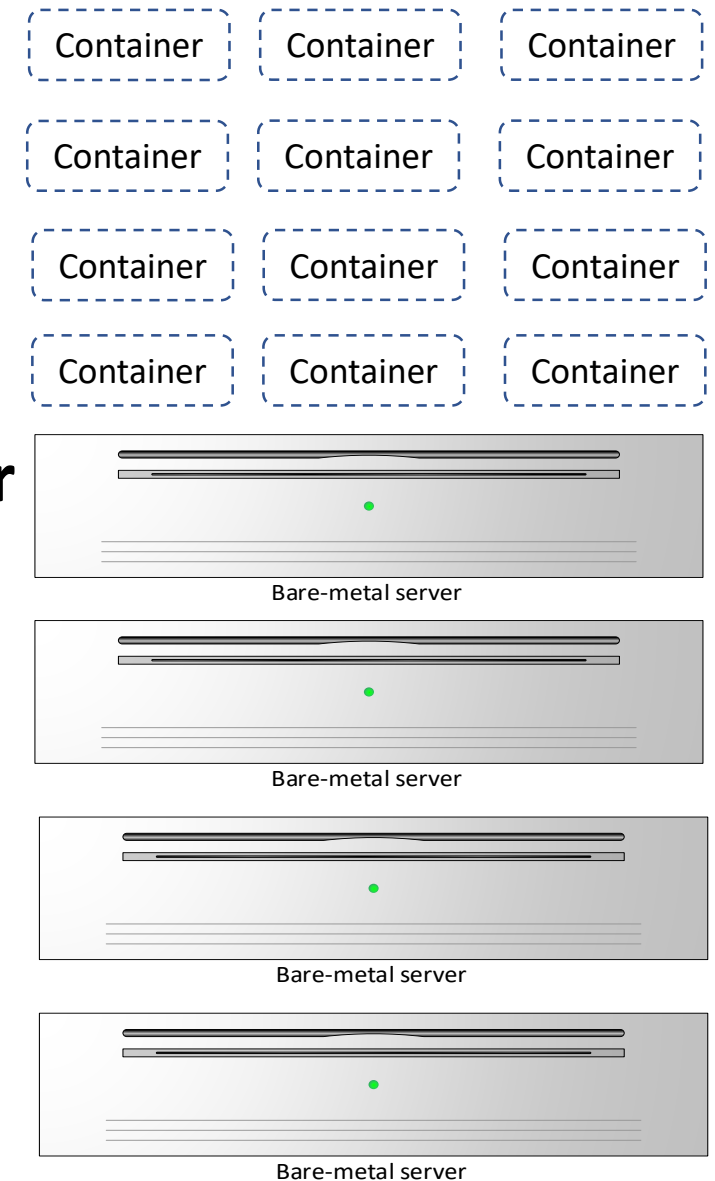


Beyond Single Node Docker

Single node



Node cluster



Multi-node production scale-out?

- 1. High-availability*
- 2. Horizontal scalability*

Beyond Manual Containers



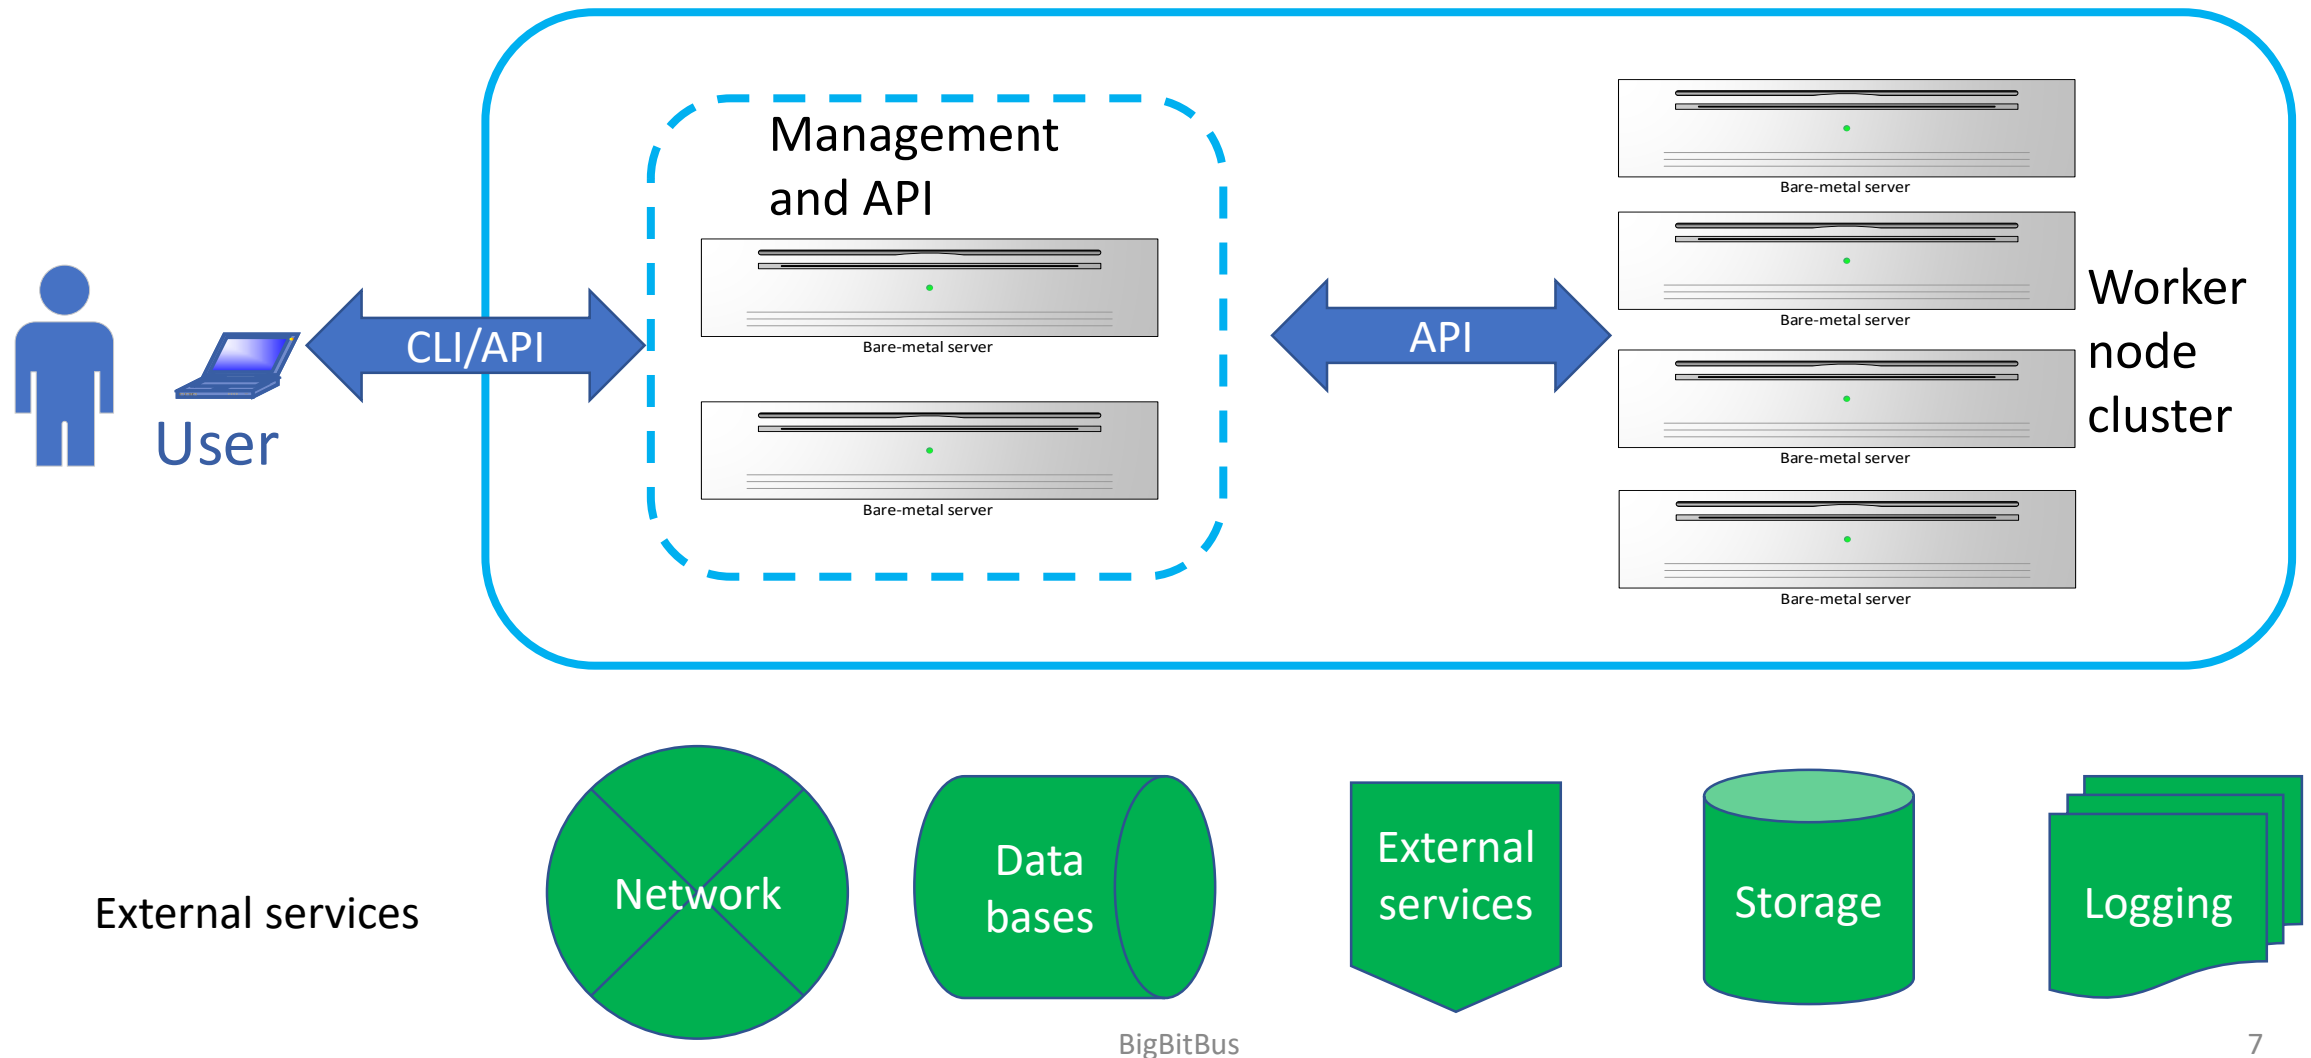
1. Multiple node container runtime cluster

- Load-balancing across multiple containers
- Service discovery
- Fairness and resource sharing among different applications

2. Production Operations Management

- Application updates & versioning
- Monitoring for failures; auto-healing
- Autonomous behavior – autoscaling and repairs
- Dealing with Failures
- Role based access controls
- Zero-downtime upgrades

Kubernetes Cluster

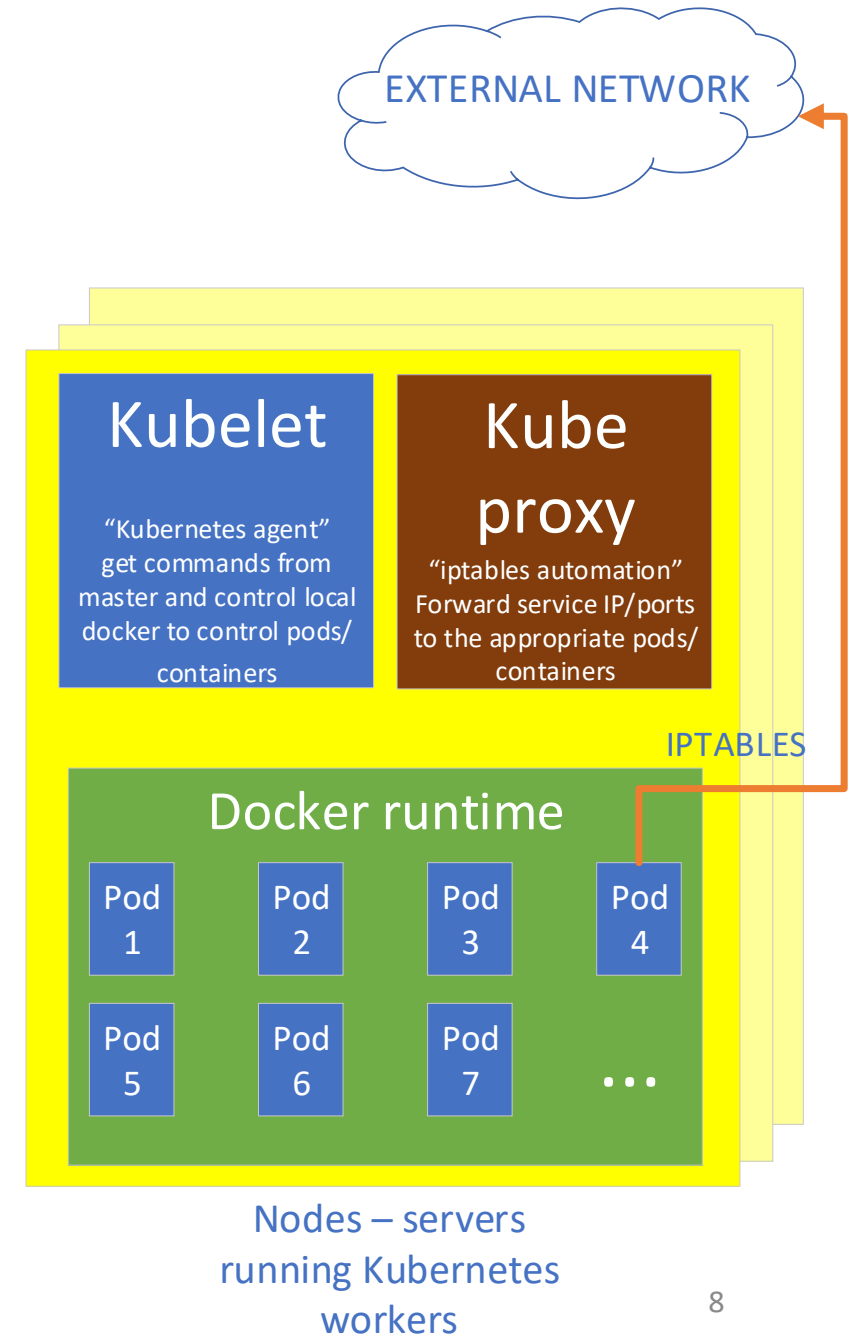



Kubernetes Architecture

Multiple components

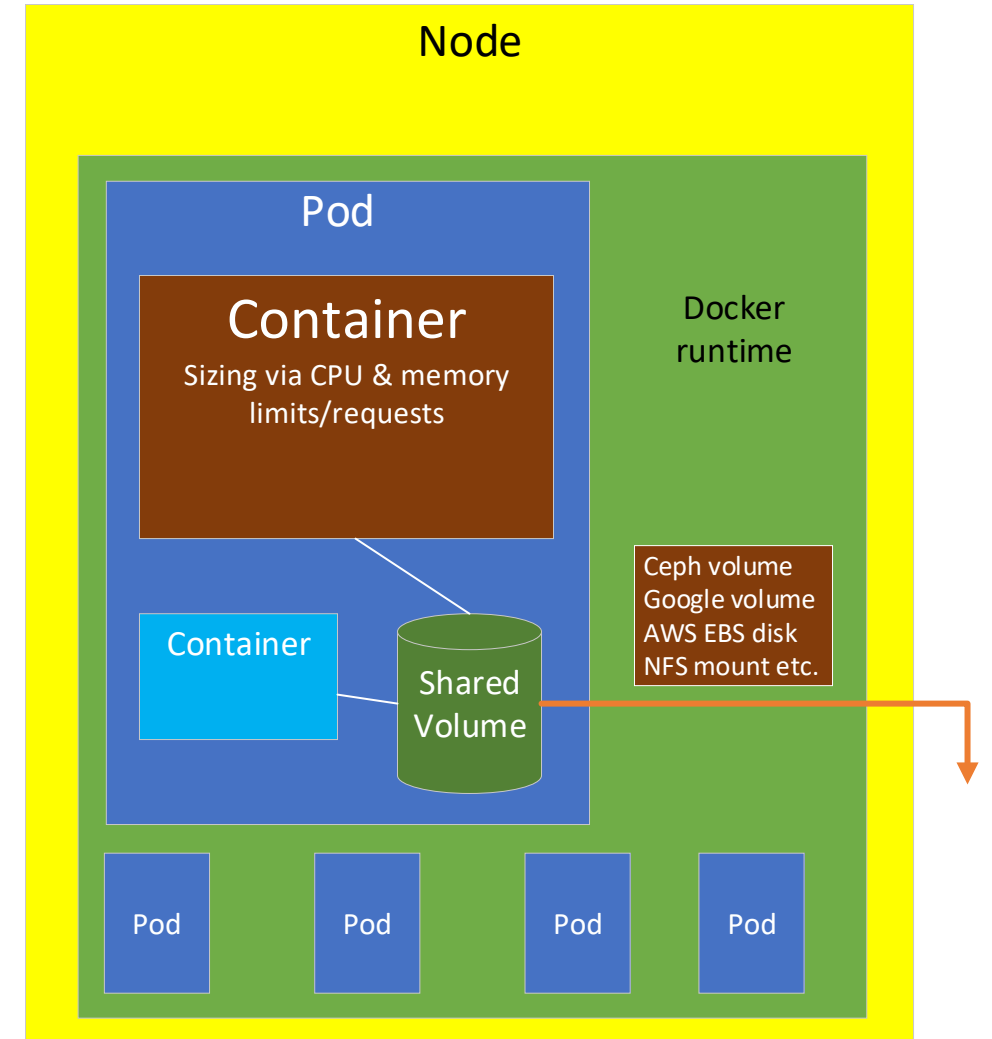

- Management plane
- Worker plane

Worker nodes can be bare-metal or virtual machines

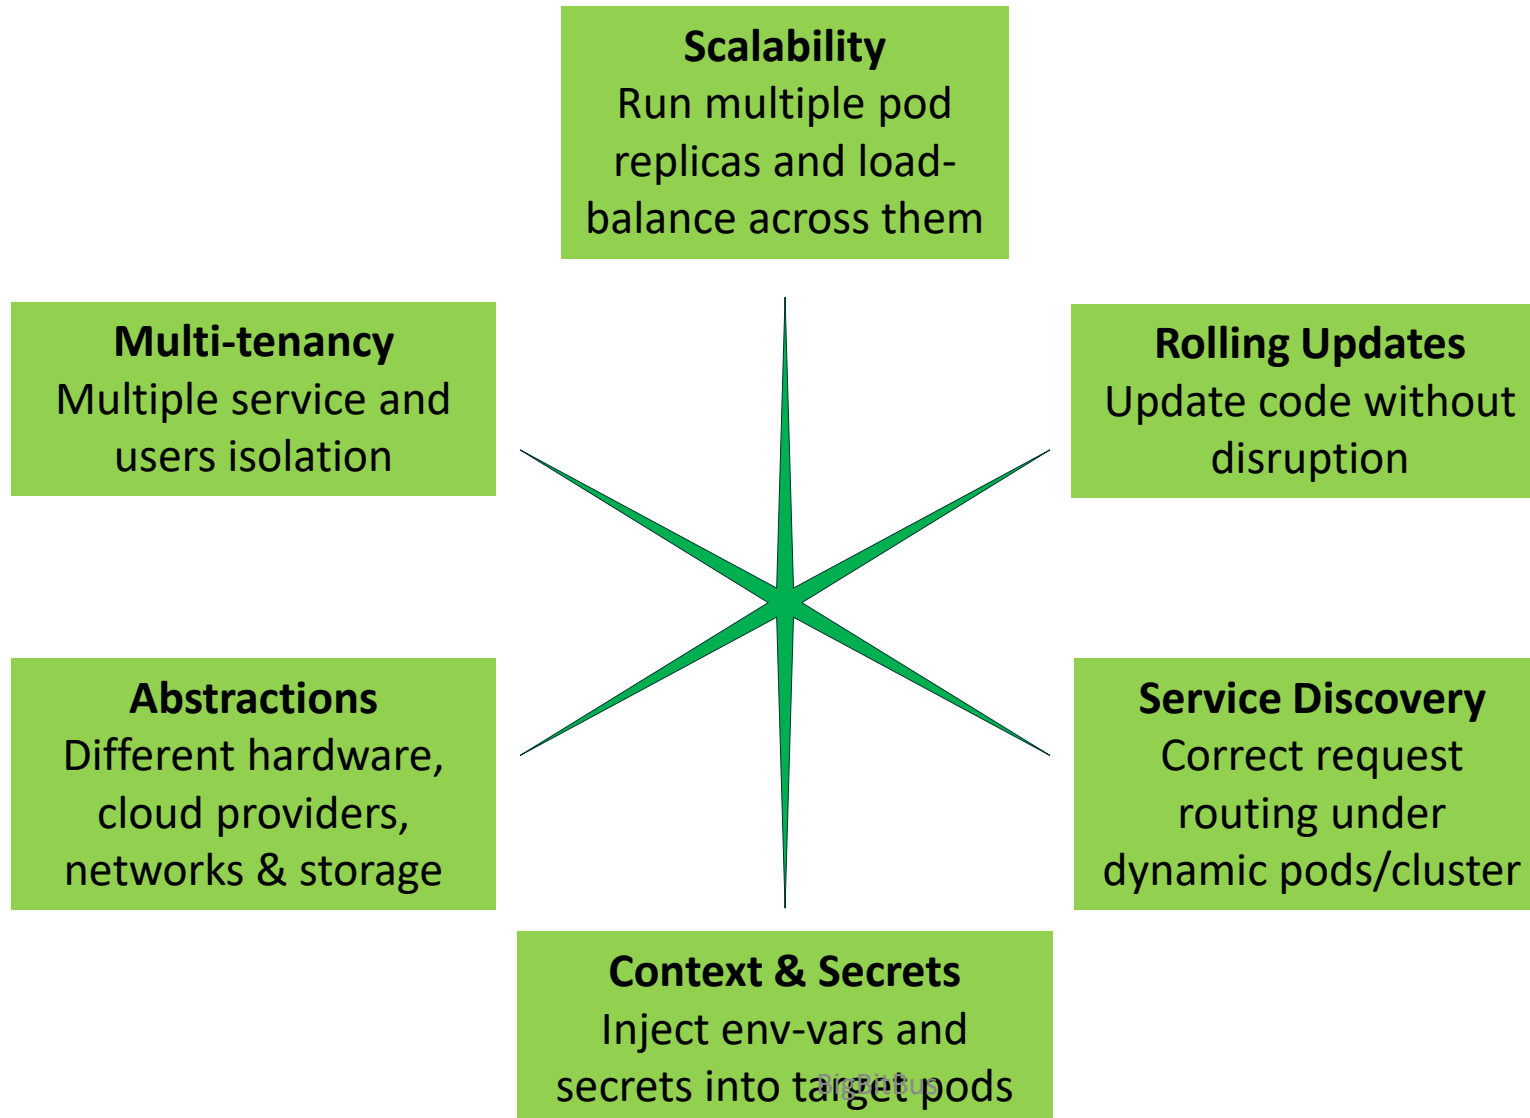


Pod

- The smallest independent “unit” of infrastructure in K8s
 - Scaling/redundancy
- A logical “host”
- Multiple containers & volumes
- Localhost and IPC connectivity
- Cannot span nodes



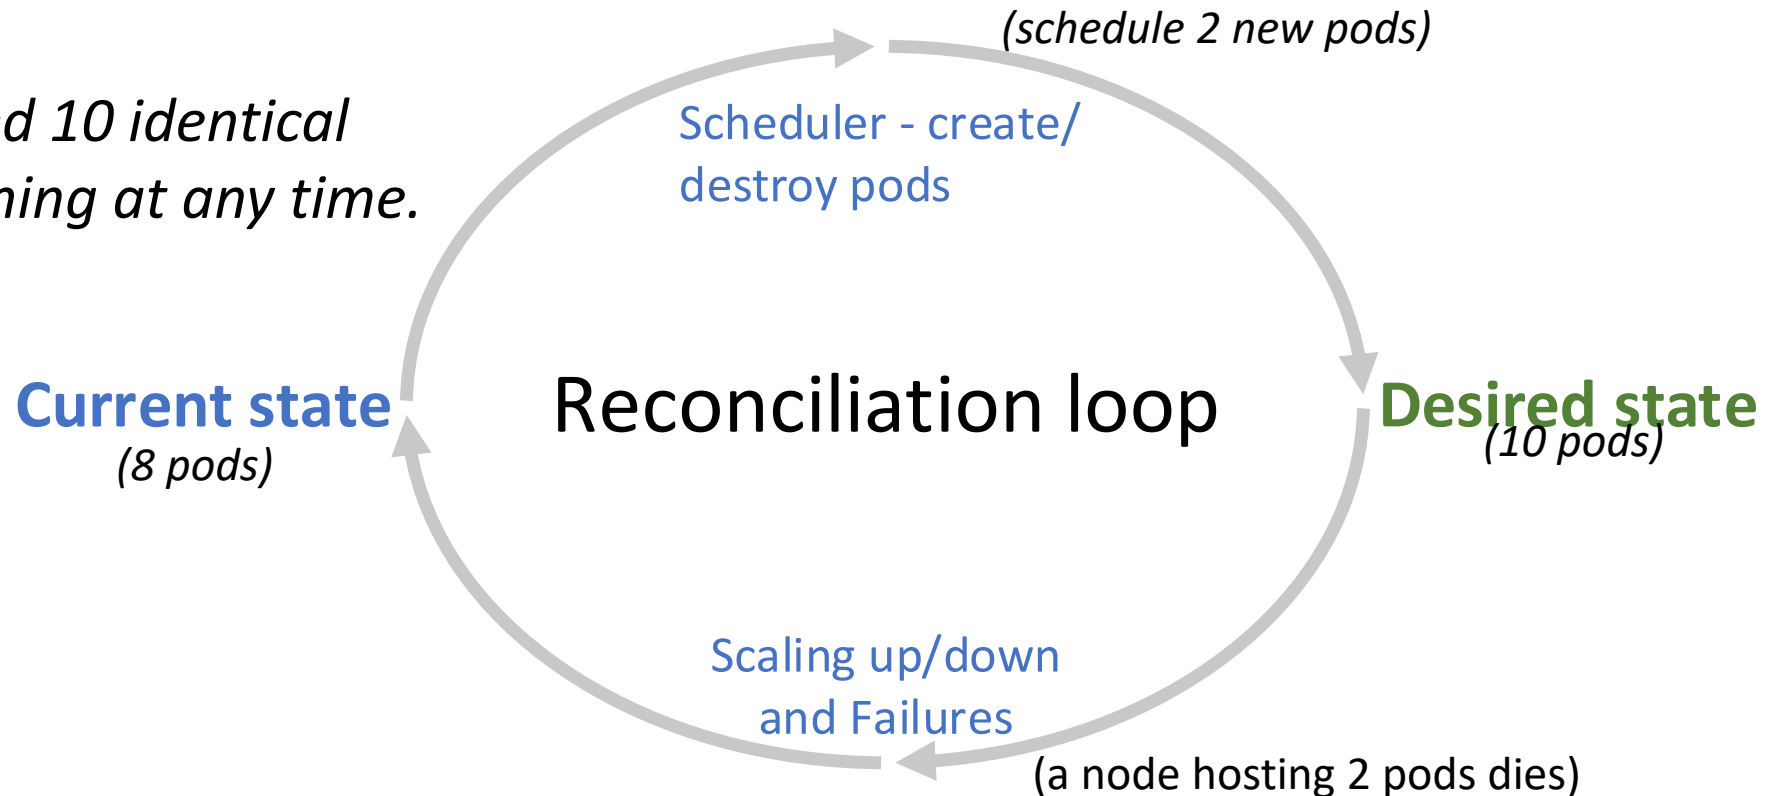
Operational Requirements



Replicaset (of pods)

- Keep a defined number of identical pod replicas running in the Kubernetes cluster

Example: We need 10 identical Django pods running at any time.



Minikube: Kubernetes on your Laptop



- Installation instructions for Windows, Mac and Linux
<https://kubernetes.io/docs/tasks/tools/install-minikube/>
- Default: 2 cores + 2 GB RAM = capable of running a few containers

Starting minikube

minikube cli


```
# start minikube, launch VM
$ minikube start
# check minikube status
$ minikube status
# start minikube GUI
$ minikube dashboard
# remember capacity constraints
$ kubectl get nodes minikube -o yaml
```

```
capacity:
  cpu: "2"
  ephemeral-storage: 16888216Ki
  hugepages-2Mi: "0"
  memory: 2038624Ki
  pods: "110"
```

BigBitBus

Dashboard

General-purpose web UI for Kubernetes clusters



Kubectl – Kubernetes CLI

kubectl --help

- Wrapper around RESTful API: GET, CREATE, DEPLOY etc. action-words

```
kubectl get nodes  
kubectl get nodes -o wide  
kubectl get node minikube -o yaml
```

More information,
yaml or json

```
kubectl config get-contexts  
kubectl config use-context aws  
kubectl get nodes
```

The same kubectl can point
to multiple k8s clusters

```
kubectl create namespace development  
kubectl get namespaces  
kubectl --namespace development get pods
```

Virtual k8s cluster via
namespaces

Run a one-off Pod

```
kubectl run my-shell --rm -i --tty --image ubuntu -- bash
```

```
ubuntu@ip-172-31-45-221:~$ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
my-shell-68974bb7f7-tbfpx          1/1     Running   0           1m
ubuntu@ip-172-31-45-221:~$ kubectl get deployments
NAME          DESIRED   CURRENT   UP-TO-DATE   AVAILABLE   AGE
my-shell      1         1         1             1           1m
ubuntu@ip-172-31-45-221:~$ kubectl delete deployment my-shell
deployment.extensions "my-shell" deleted
ubuntu@ip-172-31-45-221:~$ kubectl get pods
No resources found.
```

A pod is created, a
A deployment is also created
Cascading delete – pod also deleted

Inside the pod
Just like a “normal” Linux box

```
root@my-shell-68974bb7f7-tbfpx:/# cat /etc/hosts
# Kubernetes-managed hosts file.
127.0.0.1        localhost
::1             localhost ip6-localhost ip6-loopback
fe00::0          ip6-localnet
fe00::0          ip6-mcastprefix
fe00::1          ip6-allnodes
fe00::2          ip6-allrouters
172.17.0.5       my-shell-68974bb7f7-tbfpx
root@my-shell-68974bb7f7-tbfpx:/#
```



Running one-off pods is not the Kubernetes use-case


Deployment defined in YAML



- Deployments help manage application versions
 - Click [here](#) to see the Django deployment yaml file
 - Click [here](#) to see the Postgres deployment yaml file
- A new replicaset is created every time the deployment changes
- Deleting deployments will delete the underlying replica-sets and pods


Application updates, rollbacks, jumps to any revision

Deployments




Service Discovery

1. Assign a “cluster-IP” to a service to which other services within the cluster can connect. [Example](#)
2. Expose service externally (Load-balancer). [Example](#)
3. Connect to an external service. [Example](#)



Pods are “cattle” - expendable, replaceable - not pets

Underlying pods and hardware are ephemeral and very dynamic



Environment Variables & Secrets

- Inject meta-data into pods, e.g. based on namespace
- Configmaps for environment variables
 - Example ([development](#) vs. [production](#) namespace)
- Configmaps are very powerful e.g. for injecting custom scripts etc.
- Secrets
 - Create a secret


```
kubectl create secret generic db-password --from-literal=db-password=superSecret
```

- [Expose secret](#) inside pod as an environment variable

Rollout, Rollbacks

- TBD

Polls Application Example



- Stateless Django servers (cattle)
- All state data stored in database (pet)

Polling Application in Browser

Polls

- [What is cool about Python?](#)
- [Which cloud provider is the best?](#)
- [Which star wars movie was the best?](#)

Which cloud provider is the best?

- ☐ Digital Ocean
- ☐ Linode
- ☐ Google Cloud
- ☐ Amazon Web Services
- ☐ Microsoft Azure
- ☐ Other

Vote

State definition - models.py

```
from django.db import models

class Poll(models.Model):
    question = models.CharField(max_length=200)

    def __str__(self):
        # Python 3: def __unicode__(self):
        return self.question


class Choice(models.Model):
    poll = models.ForeignKey(Poll, on_delete=models.CASCADE)
    choice_text = models.CharField(max_length=200)
    votes = models.IntegerField(default=0)

    def __str__(self):
        # Python 3: def __unicode__(self):
        return self.choice_text
```

Stateful Databases in Kubernetes (?)


Option 1

- Kubernetes provides constructs for this (stateful sets, volumes, etc.)
- Possible cost saving (bin packing containers)
- Is self-managed state worth your developer/SRE time?



Option 2


- Use cloud-provider stateful solutions
 - AWS RDS, Google cloudsql, etc.
- Expensive, but much less operations tasks



Development and Production Setup


Development

- Helps improve development and production code parity.



Production

- Leave the stateful heavy-lifting to the experts.



[Django settings.py example](#)

[Developer workflow example](#)

Kubernetes Anti-patterns

- Large pods – avoid making pets!
- Imperative kubectl commands instead of declarative yaml
- Single namespace for multiple environments
- Lack of consistent names (images/deployments/services etc.)
- Single cluster for development and production
- Missing application-level health-checks
- Missing cluster-level (node) auto-scaling

Cost Math

Cloud provider

- Master API server [Currently \(11/2018\) \\$0.20/hour per Amazon EKS cluster \(\\$1,752 annually\)](#)
- Add worker node VMs cost
- You still need to buy all the storage/load balancers/IP addresses etc.

Minikube

- Free, but needs a good PC capable of running a fairly big VM.
- No high availability - this is only for development work

Other important concepts

- [Labels](#) and [Annotations](#): Very powerful construct to filter and track K8s resources: Required reading!
- [Horizontal pod autoscaling](#): Triggers for pod quantity adjustment
- [Statefulsets](#) (instead of replicaset): Customizing individual pods in a deployment e.g. to take on different roles.
- [Persistent volumes](#): Attach block devices to pods
- [Jobs](#) and [Cronjobs](#): Running one-off tasks or periodic pods.